

# Java Shared Memory Performance Races

CS 131 – Fall 2016

## Testing Platform

- Java (version 1.8.0\_9)
- Java(TM) SE Runtime Environment (build 1.8.0\_92-b14)
- Java HotSpot(TM) 64-Bit Server VM (build 25.92-b14, mixed mode)
- Intel(R) Xeon(R) CPU E5620 @ 2.40GHz (4 cores, 8 threads)
- 32 GB of RAM

## BetterSafe vs. Synchronized

My implementation of BetterSafe uses ReentrantLock to ensure 100% reliability. ReentrantLocks have the property of performing better under contention while retaining all of the safety of synchronized. ReentrantLocks are faster and more scalable than synchronized because of its lock polling and interruptible lock waits. Synchronized threads cannot poll, interrupt another waiting thread, or timeout after waiting for a certain period of time, while ReentrantLock can.

## BetterSorry

BetterSorry is implemented using an array of AtomicIntegers (not AtomicIntegerArray),

### 1.1. Faster than BetterSafe

BetterSorry is faster than BetterSafe because BetterSorry does not use any locking mechanism. This means that BetterSorry does not have any overhead that locks do, and that BetterSorry's threads will not wait as long as BetterSafe's threads. The increment and decrement of the race condition variables is volatile, but this does not mean it is race condition free, it means that each memory access to the variable must not be optimized away by the compiler.

### 1.2. More Reliable than Unsynchronized

BetterSorry is more reliable than Unsynchronized because it utilizes volatile behavior. Volatile guarantees that any changes made to a volatile variable are always visible to other threads. Any threads reading the variable will be able to see its most recent change, which allows BetterSorry to better avoid the infinite loop that Unsynchronized faces. This infinite loop can happen in Unsynchronized when multiple threads decrement or increment the same variable at the same time, putting the variable out of bounds (unable to increment or decrement it forever). This makes all later reads to that variable wait indefinitely.

It is possible to make BetterSorry enter the infinite loop. You can do this by setting the maxval to 1 and all of the five initial entries to 1. This forces `value[j]` to be equal to maxval,

and all threads will always return false, never allowing a swap to happen. This can also be done by setting the five initial entries to be 0.

The program to run in order to cause the infinite loop is the following:

```
$ java UnsafeMemory BetterSorry 1 1 1 1 1 1 1
```

## All Classes

### 2.1. Synchronized

Synchronized is guaranteed as a built-in keyword to be data-race free (DRF) and deadlock free. This means that only one thread can be executing swap at any given time, leading to poor scalability with large number of swaps.

### 2.2. Unsynchronized

Unsynchronized is faster than Synchronized because there are no safety overhead, but suffers from an infinite loop if any values become out of bounds. The chance of looping infinitely increases as the number of threads increases, number of swaps increases, maxval decreases, and initial values are set to numbers that are close to maxval or 0. The chance of this happening was observed to be negligible for less than 10,000 swaps, 20% for 100,000 swaps, and 60% for 1,000,000 swaps (for 2 or more threads).

### 2.3. GetNSet

GetNSet is implemented using AtomicIntegerArray, which guarantees volatile access to values. GetNSet is not DRF and is vulnerable to infinitely looping because there is no guarantee of exclusive access to the variables. The chance of looping infinitely was observed to be negligible for less than 100,000 swaps, 43% for 100,000 swaps, and 80% for 1,000,000 swaps.

### 2.4. BetterSafe

BetterSafe is implemented using ReentrantLocks, which is proven to be more scalable than Synchronized while retaining all safety characteristics. It is DRF, deadlock free, and does not ever enter the infinite loop. BetterSafe was observed to perform faster for swaps over 100,000, increasing in likelihood as the number of swaps increases (nearly 100% for 1,000,000 swaps).

### 2.5 BetterSorry

BetterSorry is implemented using an array of AtomicIntegers along with the `getAndIncrement()` and `getAndDecrement()` methods. This always performs faster than BetterSafe for swaps less than 100,000, but is slower for 1,000,000 swaps. BetterSorry is not DRF.

## Results

All tests are run with maxval as 127 and each of the 5 initial entries as 25. The reason for this is to minimize the chance of entering the infinite loop by evenly distributing the total sum over all the variables. Every combination of thread count (1, 2, 4, 8, 16, 32), swaps (100, 1,000, 10,000, 100,000, 1,000,000) is run on each class five times, and the mean of the transition time is taken for each thread count. All reported values are of unit ns/transition. Any reported values of 0 means an infinite loop was entered. If the class is not DRF, an example shell command that is likely to loop infinitely is provided below it.

		1	2	4	8	16	32
Swaps	100	32770.36	64273.66	128176.2	308381.4	819138	2699656
	1,000	5518.98	11392.28	21892.02	38377.72	107575.04	265541.8
	10,000	959.034	3046.862	5820.154	12573.54	28767.8	59960.06
	100,000	247.5492	652.95575	2195.07	0	9243.55	28688.2
	1,000,000	66.27978	0	0	0	0	0

\$ java UnsafeMemo GetNSet 32 1000000 127 25 25 25 25 25

BetterSafe		Threads					
		1	2	4	8	16	32
Swaps	100	30647.02	646248.06	167991	386853	1023911.8	2878208
	1,000	5835.436	12007.76	27957.64	57699.7	128612	316776.4
	10,000	1023.388	3027.838	6106.792	12266.74	26737.3	54843.5
	100,000	240.5926	1165.5144	2211.208	4284.948	7983.65	16856.38
	1,000,000	69.09506	627.0114	680.1482	1313.838	2762.134	5976.1

BetterSorry		Threads					
		1	2	4	8	16	32
Swaps	100	29729.42	53497	109485.12	259387	708584.2	2147244
	1,000	4772.068	8861.77	19849.18	37289.8	108005.52	256953
	10,000	793.501	2159.062	4080.088	8801.13	17237.52	38716.64
	100,000	228.5002	742.1348	1242.934	1862.938	3092.722	9462.336
	1,000,000	62.56722	248.1914	866.2644	1375.3012	3778.648	8474.21

\$ java UnsafeMemo BetterSorry 32 10000000 127 25 25 25 25 25

Based on the results, GDI should use BetterSorry, since there is some allowance for error and speed is a priority.

In order to run all of these tests efficiently, I had to write multiple bash scripts that would run test each combination of class, threads, and swaps, and then parse the results to be used in Excel to calculate average transition times.

Null		Threads					
		1	2	4	8	16	32
Swaps	100	28853.54	54522.5	121971	248350.4	763489.4	2271684
	1,000	4430.91	8239.9	18533.22	34876.14	79475	253864
	10,000	749.8884	1684.166	3780.22	7384.726	16323.94	36749.3
	100,000	207.8832	527.863	1037.2426	1313.122	2729.912	7812.028
	1,000,000	42.18246	107.8127	341.3864	1782.774	3831.254	7257.532

Synchronized		Threads					
		1	2	4	8	16	32
Swaps	100	27517.8	53589.68	114247.8	265507	789262.6	2400660
	1,000	4441.948	9010.06	16410.24	31837.8	78332.62	253864
	10,000	734.8954	1892.274	4939.464	9456.238	18867.12	40537.86
	100,000	210.1898	660.9526	2278.372	4489.976	8303.444	15483.74
	1,000,000	63.34434	380.5606	1365.652	3120.152	5316.202	8543.648

Unsynchronized		Threads					
		1	2	4	8	16	32
Swaps	100	28786.72	53579.4	123232	249653.4	809725.8	2166648
	1,000	4311.662	8785.026	16819.22	33536.98	85078.08	255960.6
	10,000	754.6728	1892.274	4939.464	9456.238	18867.12	40537.86
	100,000	213.4112	461.2558	1035.5235	1626.43	3380.17	6938.77333
	1,000,000	49.70484	0	728.078	1669.27	4783.47	9261.59

\$ java UnsafeMemory Unsynchronized 32 1000000 127 25 25 25 25 25

GetNSet	Threads
---------	---------